

Predicting Per-Unit Prices of Hass Avocados

Yunyao Zhu

3/1/2021

Introduction

In 2017, an Australian property developer claimed that millennials were spending too much money on avocado toast instead of saving for their first home (CNNMoney 2017). While the millennial avocado toast stereotype is by no means convincing, avocados are indeed a favored fruit in the U.S. According to a market report (USAID 2014), the U.S. is the world's biggest importer of avocados. As more health benefits about avocados are discovered and promoted, demand surged even more and so did the price. The goal of this project is to investigate whether we can reasonably predict the price of avocados based on past prices and geographical data.

A study (Evans 2009) investigated the various factors that might influence the avocado market. The researchers predicted that the avocado prices would likely decrease in the 2009-2010 season due to an increase in supplies and a probable decrease in demand caused by the financial crisis. Since the study was conducted over ten years ago, it seems be interesting to examine if the avocado market has changed over the years and if we can still predict the avocado prices given its supposed volatility in recent years.

Data

The dataset used in this analysis is the publicly accessible Kaggle Avocado Prices dataset, which credited the Hass Avocado Board for the collection and release of the data. This dataset contains 13 variables encompassing the per-unit prices, total volumes, regions, and sizes of Hass avocados, a cultivar of avocados. Data from the start of 2015 to the end of the first quarter of 2018 are available. Each entry represents one observation from one region in the U.S. during one week. There are over 18,000 entries in total. Details of data encoding are provided in the appendix.

To enrich this dataset, we join it with the U.S. State Demographics dataset from the `choroplethr` package (Lamstein 2021). Selected demographic data include the total population in the region, the average percentage of white and black population, the average per capita income, and the average median age.

Research Goals

Predicting the per-unit prices of Hass avocados Specifically, we want to use the data from 2015 to the first quarter of 2017 to predict the avocado prices in the following year. Doing so allows us to compare the predicted prices with the recorded prices and evaluate the accuracy of the predictions. Predicting one year of avocado prices enables us to examine any seasonal variations in pricing.

Investigating whether spatial relationships exist Does the region (where the observations were made) have any associations with the per-unit price of the avocados? If such spatial relationships are found, what might be some potential implications?

Exploratory Data Analysis

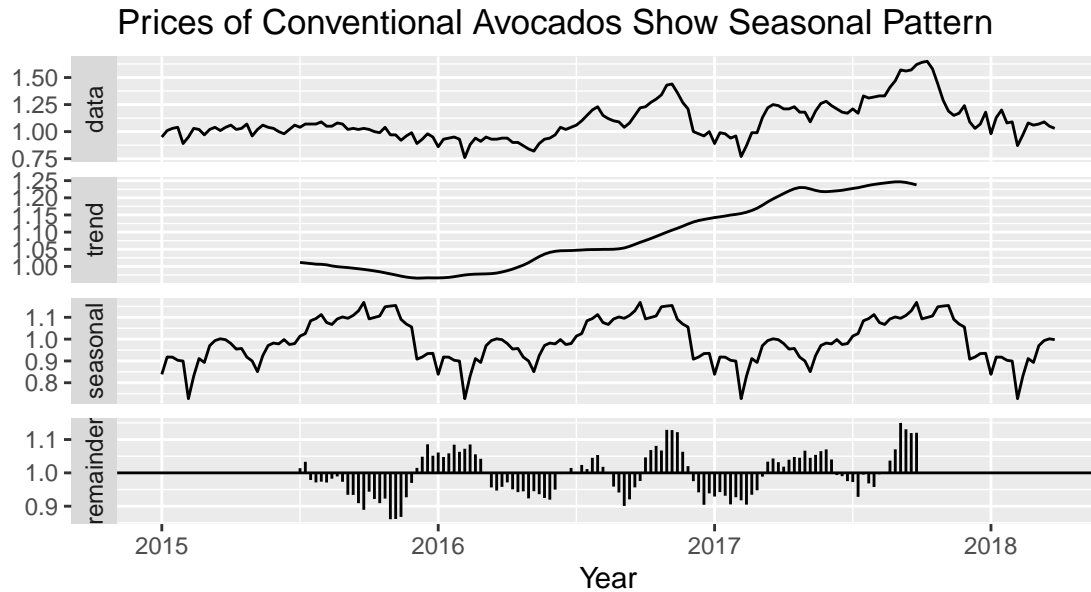


Figure 1: Time Series Plot of Per-unit Prices of Conventional Avocados in the U.S.

The per-unit price of the conventional avocados increases yearly and exhibit seasonal trends. Prices peak in the third quarter and drop to the lowest in the first quarter. Organic avocados show very similar results.

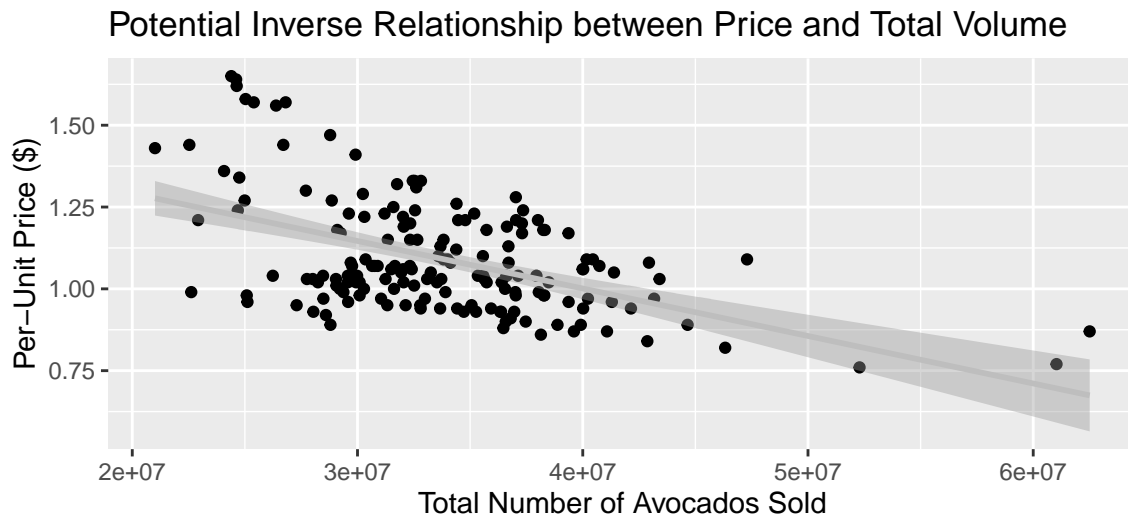


Figure 2: Scatter Plot of Per-Unit Price v.s. Total Volume of Avocados Sold

An inverse relationship seems to exist between the per-unit price and the total sales volume of the conventional avocados.

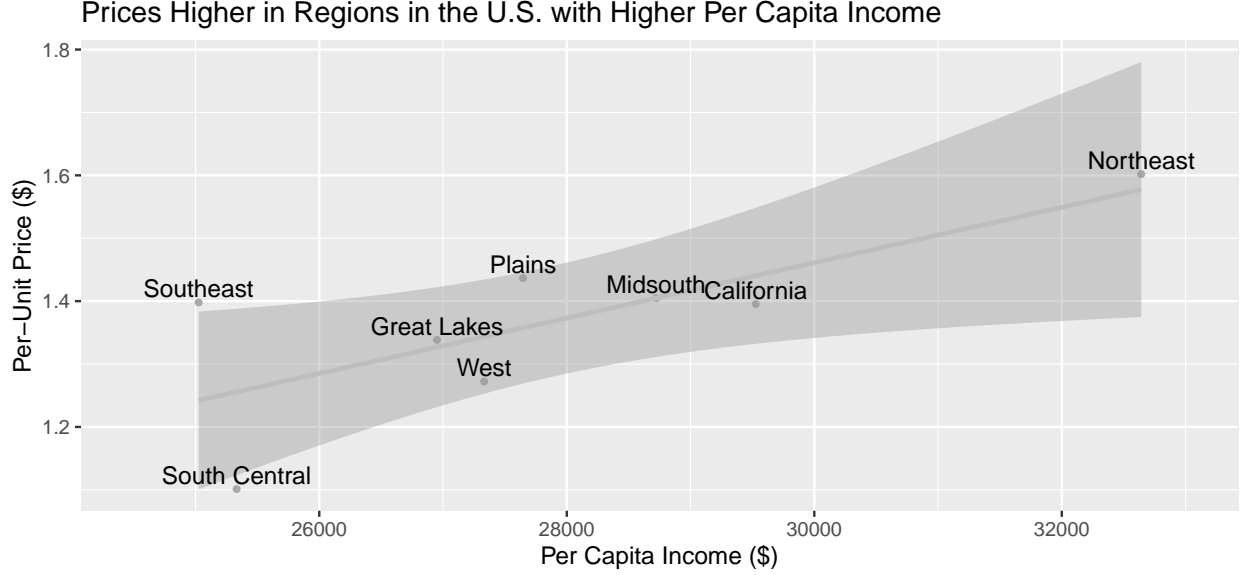


Figure 3: Scatter Plot of Per-unit Prices of Avocados v.s. Per Capita Income in Different Regions

Regions with higher per capita income seem to have higher per-unit avocado prices. An exception is the Southeast region, which has the lowest per capita income but a relatively high per-unit avocado price. Further analysis is required to better explain this observation.

Methodology

To predict the per-unit avocado prices from the first quarter of 2017 to the first quarter of 2018, we use a seasonal ARIMA with the total volume of avocados as the exogenous variable.

Model Specification

Seasonal ARIMA $(p, d, q) \times (P, D, Q)_s = (1, 0, 0) \times (0, 1, 1)_{52}$ (Rundel 2017a):

$$(1 - \phi_1 L)(1 - L^{52})y_t = \delta + (1 + \Theta_1 L)w_t$$

$$y_t - y_{t-52} = \delta + w_t + \Theta_1 w_{t-52}$$

$$y_t = \delta + y_{t-52} + w_t + \Theta_1 w_{t-52}$$

The parameter p is the number of time lags for the autoregressive (AR) model, d is the degree of differencing, and q is the order of the moving-average (MA) model. The uppercase P, D, Q denote the autoregressive, differencing, and moving average terms for the seasonal part of the ARIMA model (SAS 2021). The values of the parameters are determined using principles from the Box–Jenkins method (Rundel 2017b).

Results




Figure 4: Predicting the Weekly Per-Unit Conventional Avocado Prices from 2017-2018

The predicted per-unit prices for weeks 0-10 and weeks 35-48 are relatively close to the observed prices (differ by less than \$0.1). The predicted prices for weeks 25-35 exhibit a similar pattern to the actual prices for weeks 20-30, i.e. the predictions seem to lag by roughly 5 weeks.




Figure 5: Predicting the Monthly Per-Unit Conventional Avocado Prices from 2017-2018

Appendix

Data Encoding Details

There are 54 unique regions in the dataset. One of these regions, denoted by `TotalUS`, encompasses the avocado volumes in the country as a whole. Another 8 regions corresponds to large geographical divisions of the U.S. such as the `Northeast` region, `Southwest` region, etc. The remaining regions represent smaller geographical areas in a less consistent way, sometimes referring to a single city in a state (e.g. `Chicago`) while other times denoting cities across states (e.g. `BaltimoreWashington`). Some regions can also be ambiguous - cities such as `Albany` and `Jacksonville` can be found in more than one state. Thus, to manage the location data consistently and unambiguously, we choose to focus on the 8 geographical regions.

Seasonal ARIMA Residuals

```
##
## Ljung-Box test
##
## data: Residuals from Regression with ARIMA(1,0,0)(0,1,1)[52] errors
## Q* = 36.598, df = 7, p-value = 5.589e-06
##
## Model df: 3. Total lags used: 10
```

```
##
## Ljung-Box test
##
## data: Residuals from Regression with ARIMA(3,0,0)(0,1,1)[12] errors
## Q* = 6.3065, df = 3, p-value = 0.09762
##
## Model df: 5. Total lags used: 8
```




Figure 6: ARIMA Residuals based on Weekly Data




Figure 7: ARIMA Residuals based on Monthly Data

Bibliography

- CNNMoney. 2017. “Millionaire to Millennials: Lay Off the Avocado Toast If You Want a House.” *Cable News Network*.
- Evans, Sikavas, Edward A. & Nalampang. 2009. “Forecasting Price Trends in the U.s. Avocado (Persea Americana Mill.) Market.” *Food Distribution Research Society*, Vol. 40(2), Pages 1-10. [https://doi.org/10.1016/S0140-6736\(17\)32812-X](https://doi.org/10.1016/S0140-6736(17)32812-X).
- Lamstein, Ari. 2021. “Mapping Census Bureau Data in R with Choroplethr.” *United States Census Bureau*. <https://www.census.gov/data/academy/courses/choroplethr.html>.
- Rundel, Colin. 2017a. “Lecture 11 - Seasonal Arima.” *Duke Statistical Science Department*. http://www2.stat.duke.edu/~cr173/Sta444_Sp17/slides/Lec11.pdf.
- . 2017b. *Duke Statistical Science Department*. http://www2.stat.duke.edu/~cr173/Sta444_Sp17/slides/Lec9.pdf.
- SAS. 2021. “Notation for Arima Models.” *SAS/ETS(R) 9.3 User’s Guide*. https://support.sas.com/documentation/cdl/en/etsug/63939/HTML/default/viewer.htm#etsug_tffordet_sect016.htm.
- USAID. 2014. “The Us Market for Avocado.” *ACCESO*. https://pdf.usaid.gov/pdf_docs/PA00KP28.pdf.